

Tullgarnsvägen 2
125 71 Älvsjö
Sweden
☎ +45 44 12 18 06

✉ ivar.cashin.eriksson@gmail.com
 

Ivar Cashin Eriksson

Hello OpenAI team,

I am writing to apply for the Forward Deployed Engineer role in Stockholm. My strongest work is customer-facing delivery, owning discovery, technical scoping, system design, build, and production rollout for ambiguous customer problems. I have done that in consultancy, at Signum, and previously in Stockholm at RaySearch.

At Valcon, I worked on an invoice quality-control system for a very non-technical customer. A large part of the work was translating the business problem into an implementable ML workflow: what the model should predict, how the process should behave, what counted as an error, and how the output would fit into daily operations. The system went into production with roughly two-second response time, reached 100 percent recall within the first two weeks, and captured more than 50 business processes in the model. I care about closing the loop after deployment, measuring whether the system changes the workflow, identifying model failure modes, and feeding that evidence back into the next iteration.

That experience shaped how I think about deployment work. The technical solution only matters if it fits the workflow, survives edge cases, and gives the customer something they can actually use. I have also led data science projects across pharma, energy, IoT, and manufacturing, and managed the delivery of end-to-end production pipelines.

At Signum, I maintain a pharmaceutical forecasting platform and extend a next-best-action system inherited from a third party. In that work I use LLM embeddings for domain-specific matching workflows, including retrieval evaluation, semantic-similarity analysis, and failure analysis, and I extended the next-best-action system with explainable AI features in collaboration with UX designers. The work is practical and iterative, with attention to what people can actually use.

I have also built a full-stack computer vision and retrieval system in Project iris. It was a solo startup exploration, not a shipped customer deployment. I used YOLOv8, OpenCLIP, FastAPI, MongoDB, Qdrant, and a Chromium extension. This gave me practical full-stack experience across backend APIs, vector search, persistence, and browser-facing UX.

I have worked in Stockholm before at RaySearch, and I speak Swedish natively. I would welcome the chance to discuss how my background in customer-facing deployment, production ML, full-stack system building, and translating model behaviour into usable workflows could be useful in Stockholm.

Regards,

Ivar Cashin Eriksson.