

Ivar Cashin Eriksson

Curriculum Vitae

Øster Farimagsgade 9, 4. Th.
1353 København K
Denmark
☎ +45 44 12 18 06
✉ ivar.cashin.eriksson@gmail.com
[in](#) [Q](#)

Introduction

Senior data scientist and ML engineer building production AI systems in Databricks, with experience across forecasting, LLM embeddings, retrieval workflows, computer vision, and stakeholder-facing delivery. Based in Copenhagen and used to owning work from ambiguous business problem through implementation, productionisation, and adoption.

Selected Deployment and ML Engineering Experience

Full-Stack Data Scientist, Project Manager @ Signum Life Science

- Own and develop a Databricks-based pharmaceutical forecasting platform and LLM embedding / retrieval workflows for semantic matching, enrichment, and downstream decision support.
- Took over a next-best-action system for sales representatives and extended it with explainable AI features.

Keywords: Databricks, NLP, forecasting, embeddings, ML engineering, project management

[Read more...](#)

Principal Data Scientist @ Valcon

- Scoped and delivered an invoice quality-control ML system for a non-technical customer, translating business rules, failure modes, and daily operations into an implementable workflow.
- Production system achieved roughly two-second response time, 100% recall within the first two weeks.
- Led technical projects across many industries, including stakeholder communication and mentoring.

Keywords: production ML, stakeholder management, mentoring, explainable AI, real-time systems

[Read more...](#)

Solo Project @ Project iris

- An end-to-end AI application that detects products in e-commerce images and links similar items across a web store. Uses YOLOv8, OpenCLIP, FastAPI, MongoDB, and a custom Chromium extension.

Keywords: AI application development, startup validation, OpenCLIP, YOLOv8, vector search, FastAPI, Qdrant, computer vision

[Read more...](#)

Education

- 2015–2020 **Royal Institute of Technology, KTH – Data Science and Engineering**, GPA: 5.0 (out of 5.0)
Master: Statistical Learning and Data Analytics, Bachelor: Engineering Physics. See thesis.
Relevant coursework: Probability Theory and Statistics; Optimization; Object-Oriented Program Construction.
- 2017–2022 **Stockholm University, SU – Economics, Supplementary Bachelor's Degree**
- 2019–2019 **Swiss Federal Institute of Technology, ETH – Applied Mathematics, Exchange**

Skills and Other

- Technology Python, SQL, Databricks, MLflow, AWS, MS Azure, PyTorch, TensorFlow, Scikit-learn, SHAP, XGBoost, FastAPI, Docker, MongoDB, Qdrant, YOLO, OpenCLIP, vector search.
- Language *Native*: English, Swedish. *Intermediate*: Danish.
- Personal I enjoy working on personal projects, woodworking, climbing, bicycling, and hiking.
- Other Member of Mensa Sweden; Valcon People's Choice Award for exemplary work, leadership, team spirit, and scientific curiosity.

Detailed Work Experience

2024–**Full-Stack Data Scientist, Project Manager, Signum Life Science, Copenhagen**

Present At Signum, I maintain and develop a pharmaceutical price forecasting platform and work on LLM embedding workflows for downstream similarity search and matching. The role spans architecture, project management, data science, ML engineering, and software engineering.

I also took over a next-best-action system for sales representatives and extended it with explainable AI features in collaboration with UX designers. I support commercialisation by helping sales representatives interpret model outputs, understand explainability views, and use the system in customer-facing workflows.

In addition, I have organised an internal week-long hackathon with up to 40 participants.

2022–2024 **Principal Data Scientist, Valcon, Copenhagen**

At Valcon, I led data science projects across pharmaceuticals, energy, IoT, and manufacturing. One project involved architecting and implementing an end-to-end production pipeline for a real-time deep neural network.

Another was an invoice quality-control system that reached a response time of around two seconds in production and captured over 50 business processes. I also mentored junior colleagues, led explainable AI and Git trainings, and regularly presented technical material to non-technical stakeholders.

2021–2022 **Data Scientist, Valtech, Copenhagen**

Delivered machine learning solutions in e-commerce, including time-series forecasting with Facebook Prophet and clustering of behavioural data for personalisation. I worked closely with business stakeholders to keep outputs actionable.

2019–2021 **Researcher within Data Science, RaySearch Laboratories, Stockholm**

At RaySearch, I worked on automated radiotherapy planning using deep learning, probabilistic modelling, and multi-objective optimisation. I developed pipelines that extracted latent anatomical features from segmented 3D CT scans using autoencoders, then used learned similarity metrics and sparse Gaussian processes to support dose planning. The work fed into production treatment planning algorithms.

I collaborated closely with medical physicists and clinicians to ensure the models matched oncological standards and practical workflow requirements.

Projects

Project iris End-to-end image-processing and retrieval system for e-commerce imagery. Built object detection, semantic embeddings, vector indexing, similarity search, backend APIs, persistence, dynamic scraping, and a Chromium extension using YOLOv8, OpenCLIP, FastAPI, MongoDB, and Qdrant.

Explored through H&M's startup ecosystem, including customer discovery, product scoping, and market validation.

resume_builder3 Source-grounded agentic workflow for job applications. The system uses structured profile YAML, job listings, LaTeX templates, and previous examples to generate tailored CVs and cover letters. It also produces evidence mappings and gap reports so major claims can be traced back to source material and weak fit is visible before submission. I extended the workflow with a human-in-the-loop job-hunter agent that crawls the web for relevant adverts, sends a daily digest, waits for my response, creates GitHub issues for selected roles, triggers GitHub Actions/Codex to generate a branch with listing information and application materials, and opens a PR for review. This application is of course generated with this system!